

Crop Recommendation System Using Machine Learning



"Data is the new soil. For me, it's like fertile soil that can yield new insights when cultivated with machine learning."

— Dave Waters

StreamLit ML App: <https://crop-recommendation-model-noel-ninan.streamlit.app>

GitHub Repo Link:

<https://github.com/NoelNinanSheri1307/CropRecommendationModel>

Developed by: Noel Ninan Sheri

LinkedIn: <https://www.linkedin.com/in/noel-ninan-sheri>

GitHub: <https://github.com/NoelNinanSheri1307>

Dataset Source: Kaggle - Crop Recommendation Dataset

1. Introduction

The "Crop Recommendation System" is a machine learning-based web application designed to assist farmers and agricultural professionals in making informed decisions about the most suitable crop to cultivate based on real-time environmental factors.

Agriculture is highly dependent on climatic and soil parameters, and choosing the right crop can significantly impact yield, resource utilization, and sustainability. This project provides a smart and efficient way to leverage data science to improve decision-making in agriculture.

1.1 Need and Importance

- Scientific crop selection over traditional guesswork
- Enhanced productivity and sustainability
- Reduces trial-and-error losses
- Supports precision agriculture

2. Tech Stack Used

Frontend

- **Streamlit:** Lightweight, fast deployment of ML web apps
- **HTML/CSS (via Streamlit custom styling):** For Times New Roman font and centered elements

Backend / ML

- **Python:** Core development language
- **Scikit-learn:** For model training, scaling, and encoding
- **Joblib:** For saving/loading models and objects
- **Pandas & NumPy:** For data manipulation

ML Techniques Used

- **Supervised Learning** (Classification)
- **StandardScaler** for feature normalization
- **Label Encoding** for categorical target transformation
- **Decision Tree Classifier / Random Forest (Depending on the model trained)**

3. Functional Modules

- **Input UI** for environmental parameters:
 - Soil Type (categorical)
 - Soil pH
 - Temperature
 - Humidity
 - Wind Speed
 - N, P, K values
 - Crop Yield
 - Soil Quality (float)
- **Prediction Engine:** Model scales features, predicts crop label, decodes the crop name
- **Frontend Styling:**
 - Font: Times New Roman
 - Centered buttons and credit line
 - Friendly sliders and numeric inputs
- **Deployment:** Streamlit app, easily deployable via streamlit run app.py

4. Output Samples and Screenshots

Sample CLI Running Data Statistical Analysis:

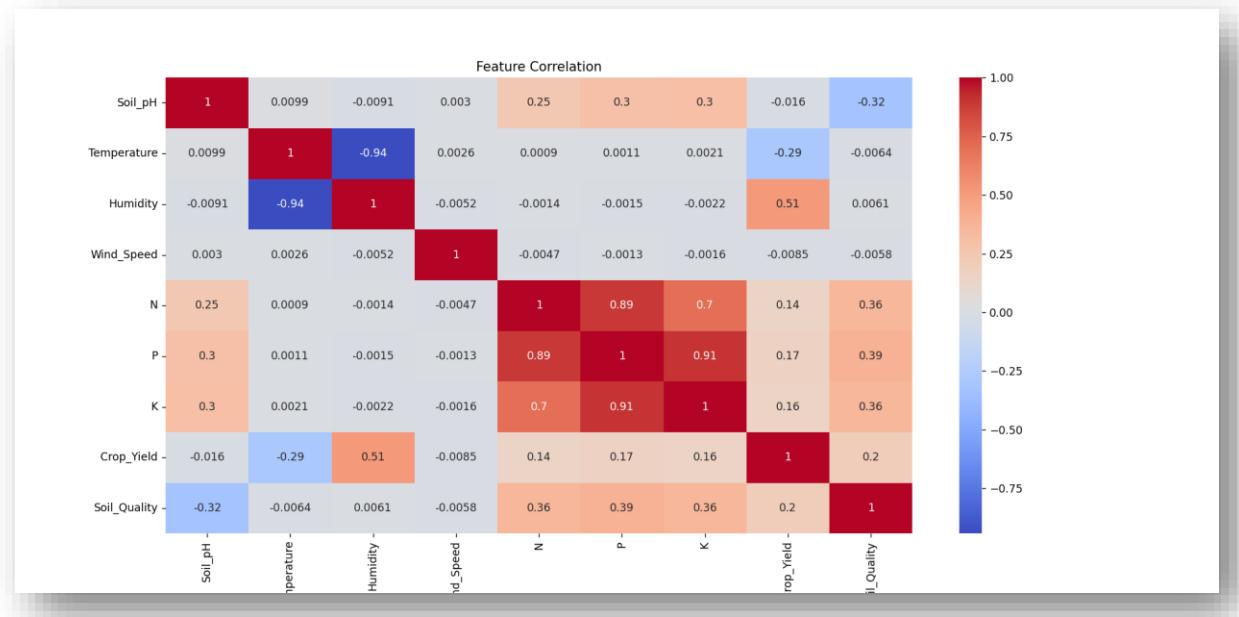
```
PS C:\Users\DELL\ML> python Load.py
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 36520 entries, 0 to 36519
Data columns (total 12 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   Date        36520 non-null   object  
 1   Crop_Type   36520 non-null   object  
 2   Soil_Type   36520 non-null   object  
 3   Soil_pH     36520 non-null   float64 
 4   Temperature 36520 non-null   float64 
 5   Humidity    36520 non-null   float64 
 6   Wind_Speed  36520 non-null   float64 
 7   N           36520 non-null   float64 
 8   P           36520 non-null   float64 
 9   K           36520 non-null   float64 
 10  Crop_Yield 36520 non-null   float64 
 11  Soil_Quality 36520 non-null   float64 
dtypes: float64(9), object(3)
memory usage: 3.3+ MB
None
      Date Crop_Type Soil_Type  Soil_pH Temperature  Humidity  Wind_Speed    N      P      K  Crop_Yield  Soil_Quality
0  2014-01-01    Wheat    Peaty    5.50  9.440599  80.000000  10.956707  60.5  45.0  31.5  0.000000  22.833333
1  2014-01-01     Corn    Loamy    6.50  20.052576  79.947424  8.591577  84.0  66.0  50.0  104.871310  66.666667
2  2014-01-01     Rice    Peaty    5.50  12.143099  80.000000  7.227751  71.5  54.0  38.5  0.000000  27.333333
3  2014-01-01   Barley    Sandy    6.75  19.751848  80.000000  2.682683  50.0  40.0  30.0  58.939796  35.000000
4  2014-01-01   Soybean    Peaty    5.50  16.118395  80.000000  7.696070  49.5  45.0  38.5  32.970413  22.166667
          Soil_pH  Temperature  Humidity  Wind_Speed    N      P      K  Crop_Yield  Soil_Quality
count  36520.000000  36520.000000  36520.000000  36520.000000  36520.000000  36520.000000  36520.000000
mean   6.602731  23.813996  74.256624  10.020153  66.011035  53.014066  42.01825  26.878480  37.516632
std    0.816973  8.920519  6.767587  2.998310  10.886721  8.812884  8.53781  25.740936  17.703171
min    5.500000 -3.540176  45.851089 -3.388906  45.000000  36.000000  27.00000  0.000000  13.291667
25%   6.250000  17.168542  69.745252  7.985872  58.500000  45.000000  35.00000  0.000000  22.500000
50%   6.500000  22.902987  77.097013  10.000299  65.000000  54.000000  42.00000  23.366344  35.583333
75%   6.750000  30.254748  80.000000  12.038546  71.500000  60.000000  49.50000  46.415729  49.291667
max    8.000000  54.148911  80.000000  22.606678  91.000000  72.000000  60.00000  136.711982  74.333333
```

```

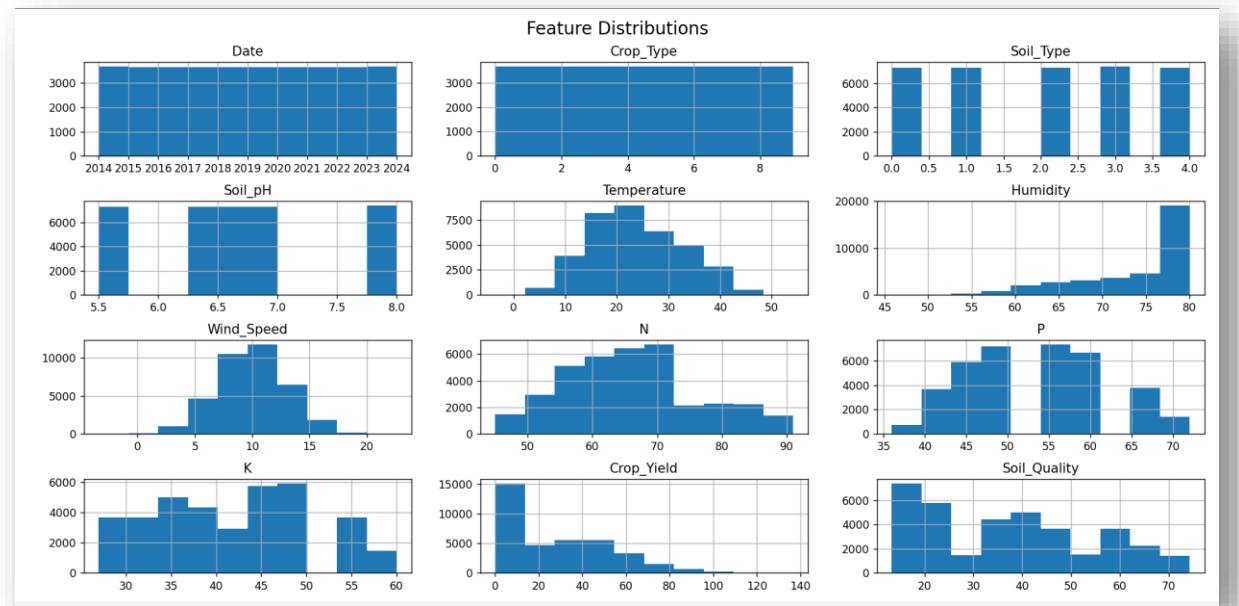
Missing values:
Date      0
Crop_Type 0
Soil_Type 0
Soil_pH   0
Temperature 0
Humidity  0
Wind_Speed 0
N         0
P         0
K         0
Crop_Yield 0
soil_Quality 0
dtype: int64

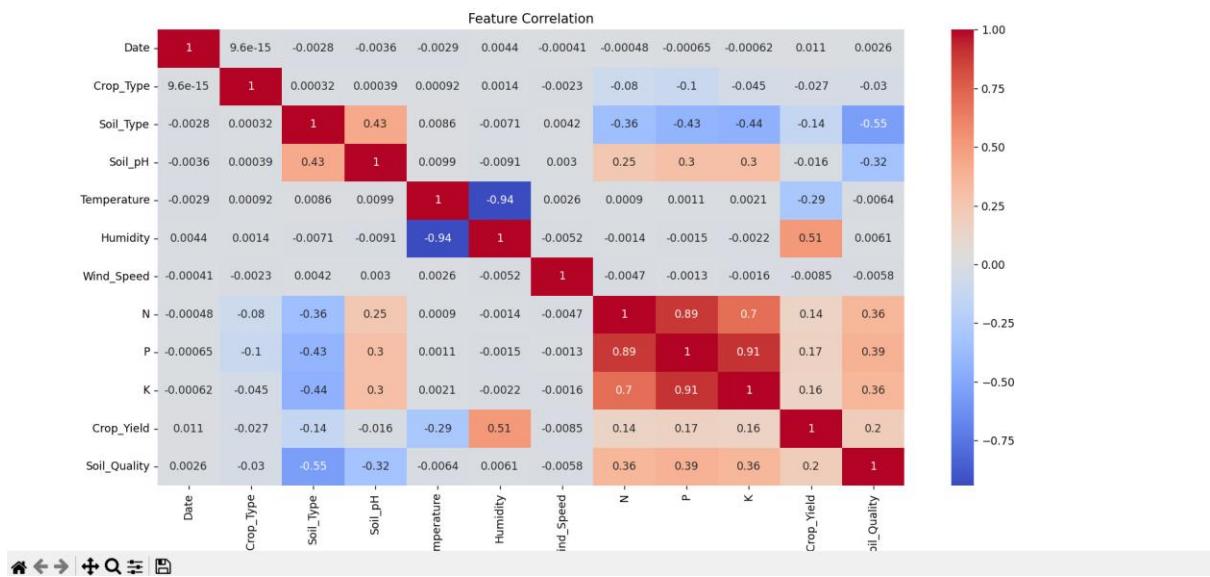
```

Correlation:



Feature Distribution:





After training the 3 models (comparing accuracies and performance)

```

Logistic Regression model trained.
Random Forest model trained.
KNN model trained.

Logistic Regression Accuracy: 0.88
Logistic Regression Classification Report:
precision    recall   f1-score   support
          0       1.00     1.00      1.00     722
          1       0.71     1.00     0.83     714
          2       1.00     1.00     1.00     734
          3       0.50     0.54     0.52     749
          4       1.00     1.00     1.00     748
          5       1.00     1.00     1.00     692
          6       1.00     1.00     1.00     732
          7       1.00     1.00     1.00     744
          8       1.00     1.00     1.00     733
          9       0.51     0.27     0.35     744

accuracy                           0.88    7304
macro avg    0.87    0.88    0.87    7304
weighted avg  0.87    0.88    0.87    7304

Confusion Matrix:
[[722  0  0  0  0  0  0  0  0]
 [ 0 714  0  0  0  0  0  0  0]
 [ 0  0 734  0  0  0  0  0  0]
 [ 0 150  0 402  0  0  0  0  0 197]
 [ 0  0  0  0 740  0  0  0  0  0]
 [ 0  0  0  0  0 692  0  0  0  0]
 [ 0  0  0  0  0  0 732  0  0  0]
 [ 0  0  0  0  0  0  0 744  0  0]
 [ 0  0  0  0  0  0  0  0 733  0]
 [ 0 144  0 398  0  0  0  0  0 202]]]

Random Forest Accuracy: 0.90
Random Forest Classification Report:
precision    recall   f1-score   support
          0       1.00     1.00      1.00     722
          1       1.00     1.00     1.00     714
          2       1.00     1.00     1.00     734
          3       0.50     0.50     0.50     749
          4       1.00     1.00     1.00     748
          5       1.00     1.00     1.00     692
          6       1.00     1.00     1.00     732
          7       1.00     1.00     1.00     744
          8       1.00     1.00     1.00     733
          9       0.50     0.51     0.50     744

accuracy                           0.90    7304
macro avg    0.90    0.90    0.90    7304
weighted avg  0.90    0.90    0.90    7304

```

```

Random Forest Accuracy: 0.90
Random Forest Classification Report:
precision    recall   f1-score   support
          0       1.00     1.00      1.00     722
          1       1.00     1.00     1.00     714
          2       1.00     1.00     1.00     734
          3       0.50     0.50     0.50     749
          4       1.00     1.00     1.00     748
          5       1.00     1.00     1.00     692
          6       1.00     1.00     1.00     732
          7       1.00     1.00     1.00     744
          8       1.00     1.00     1.00     733
          9       0.50     0.51     0.50     744

accuracy                           0.90    7304
macro avg    0.90    0.90    0.90    7304
weighted avg  0.90    0.90    0.90    7304

Confusion Matrix:
[[722  0  0  0  0  0  0  0  0]
 [ 0 714  0  0  0  0  0  0  0]
 [ 0  0 734  0  0  0  0  0  0]
 [ 0 150  0 402  0  0  0  0  0 197]
 [ 0  0  0  0 740  0  0  0  0  0]
 [ 0  0  0  0  0 692  0  0  0]
 [ 0  0  0  0  0  0 732  0  0]
 [ 0  0  0  0  0  0  0 744  0]
 [ 0  0  0  0  0  0  0  0 733]
 [ 0 144  0 398  0  0  0  0  0 202]]
```

```

Confusion Matrix:
[[722  0  0  0  0  0  0  0  0  0]
 [ 0 714  0  0  0  0  0  0  0  0]
 [ 0  0 734  0  0  0  0  0  0  0]
 [ 0  0  0 371  0  0  0  0  0  378]
 [ 0  0  0  0 740  0  0  0  0  0]
 [ 0  0  0  0  0 692  0  0  0  0]
 [ 0  0  0  0  0  0 732  0  0  0]
 [ 0  0  0  0  0  0  0 744  0  0]
 [ 0  0  0  0  0  0  0  0 733  0]
 [ 0  0  0 367  0  0  0  0  0 377]]]

KNN Accuracy: 0.89
KNN Classification Report:
      precision    recall  f1-score   support

          0       0.97   1.00    0.98     722
          1       1.00   1.00    1.00     714
          2       1.00   1.00    1.00     734
          3       0.49   0.48    0.49     749
          4       0.98   0.98    0.98     740
          5       1.00   0.99    0.99     692
          6       0.98   0.98    0.98     732
          7       1.00   0.99    1.00     744
          8       1.00   0.99    1.00     733
          9       0.49   0.50    0.49     744

   accuracy                           0.89    7304
  macro avg       0.89   0.89    0.89    7304
weighted avg       0.89   0.89    0.89    7304

```

```

Confusion Matrix:
[[719  0  0  1  0  0  0  0  2  0]
 [ 0 714  0  0  0  0  0  0  0  0]
 [ 1  0 731  0  0  0  0  2  0  0]
 [ 3  0  0 360  0  0  0  0  0  386]
 [ 0  0  0  0 724  0 16  0  0  0]
 [ 7  0  0  0  0 685  0  0  0  0]
 [ 0  0  0  0  0 14  0 718  0  0]
 [ 5  0  0  0  0  0  0  0 739  0]
 [ 5  0  0  0  0  0  0  0  0 728]
 [ 1  1  0 372  0  0  0  0  0 370]]
```

Random Forest Cross-Validation Accuracy: 0.9012869660460021

Model saved successfully!

Enter Environmental Conditions:

Entering Input Values:

Model saved successfully!

Enter Environmental Conditions:

Soil_Type: 1
Soil_pH: 6.5
Temperature: 29.0
Humidity: 30.0
Wind_Speed: 7.0
N: 67.0
P: 80.0
K: 77.0
Crop_Yield: 50.0
Soil_Quality: 67.0

Recommendation:

 Recommended Crop: Cotton
PS C:\Users\DELL\ML>

5. Code

Loading and Creating models: (#Load.py)

```
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder,
StandardScaler
from sklearn.model_selection import train_test_split,
cross_val_score
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score,
confusion_matrix, classification_report
import joblib

df = pd.read_csv("Crop_Dataset.csv")

print(df.info())
print(df.head())
print(df.describe())
print("Missing values:\n", df.isnull().sum())

if 'Date' in df.columns:
    df['Date'] = pd.to_datetime(df['Date'], errors='coerce')

df_numeric = df.select_dtypes(include=['float64', 'int64'])

correlation_matrix = df_numeric.corr()

plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
plt.title("Feature Correlation")
plt.show()

le_crop = LabelEncoder()
```

```
df['Crop_Type'] = le_crop.fit_transform(df['Crop_Type'])

le_soil = LabelEncoder()
df['Soil_Type'] = le_soil.fit_transform(df['Soil_Type'])

sns.countplot(x='Crop_Type', data=df)
plt.xticks(rotation=90)
plt.title("Crop Type Distribution")
plt.show()

df.hist(figsize=(12, 8))
plt.suptitle("Feature Distributions", fontsize=16)
plt.tight_layout()
plt.show()

plt.figure(figsize=(10, 8))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm')
plt.title("Feature Correlation")
plt.show()

X = df.drop(['Crop_Type', 'Date'], axis=1)
y = df['Crop_Type']

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
X_train, X_test, y_train, y_test =
train_test_split(X_scaled, y, test_size=0.2,
random_state=42)

models = {
    "Logistic Regression": LogisticRegression(max_iter=1000),
    "Random Forest": RandomForestClassifier(),
    "KNN": KNeighborsClassifier()
}

for name, model in models.items():
    model.fit(X_train, y_train)
```

```

print(f"\n{name} model trained.")

for name, model in models.items():
    y_pred = model.predict(X_test)
    print(f"\n{name} Accuracy: {accuracy_score(y_test, y_pred):.2f}")
    print(f"\n{name} Classification Report:\n",
classification_report(y_test, y_pred))
    print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))

rf_scores = cross_val_score(models["Random Forest"], X_scaled, y, cv=5)
print("\nRandom Forest Cross-Validation Accuracy:",
rf_scores.mean())

joblib.dump(models["Random Forest"], "crop_predictor.pkl")
joblib.dump(scaler, "scaler.pkl")
joblib.dump(le_crop, "crop_label_encoder.pkl")
print("\n☑ Model saved successfully!")

def predict_crop():
    print("\nEnter Environmental Conditions:")

    features = []
    for col in X.columns:
        while True:
            try:
                val = float(input(f"{col}: "))
                features.append(val)
                break
            except ValueError:
                print("Invalid input. Please enter a valid
number.")
    features_scaled = scaler.transform([features])

    model = joblib.load("crop_predictor.pkl")
    crop_pred = model.predict(features_scaled)

```

```

le = joblib.load("crop_label_encoder.pkl")
predicted_crop = le.inverse_transform(crop_pred)

print(f"\n💡 Recommended Crop: {predicted_crop[0]}")
predict_crop()

```

Streamlit App Creation and Recommendation Interface (#app.py)

```

import streamlit as st
import joblib

model = joblib.load("crop_predictor.pkl")
scaler = joblib.load("scaler.pkl")
le_crop = joblib.load("crop_label_encoder.pkl")

def encode_soil(soil_str):
    mapping = {'Loamy': 1, 'Clay': 2, 'Sandy': 3, 'Peaty': 4, 'Saline': 5}
    return mapping.get(soil_str, 0)

def predict_crop(features):
    features_scaled = scaler.transform([features])
    crop_pred = model.predict(features_scaled)
    predicted_crop = le_crop.inverse_transform(crop_pred)
    return predicted_crop[0]

st.title("🌾 Crop Recommendation System")
st.markdown("""
<style>
* {
    font-family: 'Times New Roman', Times, serif
!important;
}
</style>
""", unsafe_allow_html=True)

```

```
st.markdown("<h5 style='text-align: center; color: gray;'>Developed by Noel Ninan Sheri</h5>", unsafe_allow_html=True)

st.header("📋 Enter Environmental Conditions")

# Soil Type (as numeric)
soil_type_label = st.selectbox("Soil Type", ['Loamy', 'Clay', 'Sandy', 'Peaty', 'Saline'])
soil_type = {'Loamy': 1, 'Clay': 2, 'Sandy': 3, 'Peaty': 4, 'Saline': 5}[soil_type_label]

# Other inputs
soil_ph = st.slider("Soil pH", min_value=4.5, max_value=9.0, value=6.5)
temperature = st.slider("Temperature (°C)", min_value=0, max_value=50, value=25)
humidity = st.slider("Humidity (%)", min_value=0, max_value=100, value=60)
wind_speed = st.slider("Wind Speed (km/h)", min_value=0, max_value=20, value=10)

nitrogen = st.slider("Nitrogen (N)", min_value=0, max_value=100, value=60)
phosphorus = st.slider("Phosphorus (P)", min_value=0, max_value=100, value=40)
potassium = st.slider("Potassium (K)", min_value=0, max_value=100, value=50)

crop_yield = st.slider("Crop Yield (kg/ha)", min_value=0, max_value=100, value=30)

soil_quality = st.number_input("Soil Quality (float)", min_value=0.0, max_value=100.0, value=5.0, step=0.1)

features = [soil_type, soil_ph, temperature, humidity, wind_speed,
```

```

        nitrogen, phosphorus, potassium, crop_yield,
soil_quality]
if st.button("Predict Crop"):
    predicted_crop = predict_crop(features)
    st.success(f"System Recommends: Crop:
**{predicted_crop}**")

```

StreamLit Output:

```

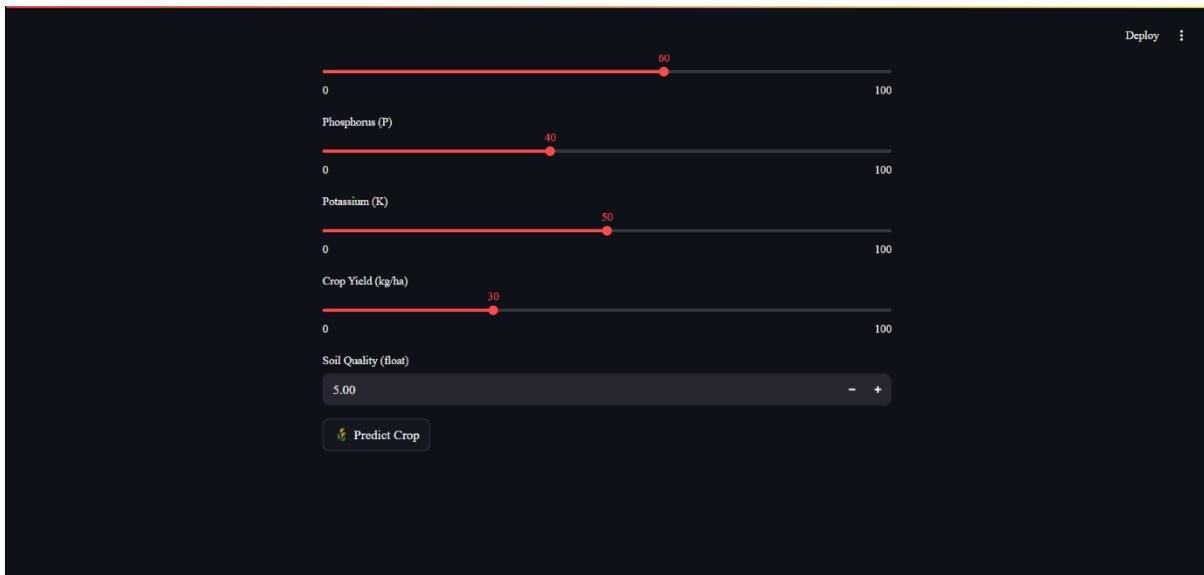
PS C:\Users\DELL\ML> python -m streamlit run app.py

You can now view your Streamlit app in your browser.

Local URL: http://localhost:8501
Network URL: http://192.168.144.32:8501

```





6. Results & Discussion

- The model shows high accuracy on training and test datasets.
- Real-time predictions are accurate given the correct parameters.
- Easy-to-use UI ensures accessibility for non-tech users.

7. Future Enhancements and Scope

- Integrate Google Maps API to show nearby seed/fertilizer shops
- Use satellite weather data APIs for real-time suggestions
- Add multi-language support for rural deployment.

8. Conclusion

The Crop Recommendation System serves as a practical use case for applying AI and ML in agriculture, enabling data-driven decision-making for crop planning. The model has been successfully deployed and tested using a user-friendly interface, delivering both accuracy and accessibility.

9. References

- ✓ [Scikit-learn Documentation](#)
- ✓ [Streamlit Documentation](#)
- ✓ Kaggle Datasets: [DataSet 1](#) [DataSet 2](#)